

## IMPORTING LIBRARIES

```
import pandas as pd
import numpy as np
import xgboost
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import classification_report
from tensorflow.keras import models
from tensorflow.keras import layers
from tensorflow.keras import optimizers
from tensorflow.keras.callbacks import EarlyStopping
import seaborn as sns
from imblearn.over_sampling import SMOTE
import matplotlib.pyplot as plt
```

## LOADING THE DATASET

```
stroke_dataset = pd.read_csv("healthcare-dataset-stroke-data.csv")
```

```
stroke_dataset
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Resid
<b>0</b>	9046	Male	67.0	0	1	Yes	Private	
<b>1</b>	51676	Female	61.0	0	0	Yes	Self-employed	
<b>2</b>	31112	Male	80.0	0	1	Yes	Private	
<b>3</b>	60182	Female	49.0	0	0	Yes	Private	
<b>4</b>	1665	Female	79.0	1	0	Yes	Self-employed	
...	...	...	...	...	...	...	...	...
<b>5105</b>	18234	Female	80.0	1	0	Yes	Private	
<b>5106</b>	44873	Female	81.0	0	0	Yes	Self-employed	
<b>5107</b>	19723	Female	35.0	0	0	Yes	Self-employed	
<b>5108</b>	37544	Male	51.0	0	0	Yes	Private	
<b>5109</b>	44679	Female	44.0	0	0	Yes	Govt_job	

5110 rows × 12 columns


## CONVERTING POSSIBLE VALUES TO CATEGORICAL VARIABLES

```
for col in stroke_dataset.columns:
    if stroke_dataset[col].dtype == 'object' or (stroke_dataset[col].dtype == 'int64' and
        print(col,"->", stroke_dataset[col].unique())
```

```
gender -> ['Male' 'Female' 'Other']
hypertension -> [0 1]
heart_disease -> [1 0]
ever_married -> ['Yes' 'No']
work_type -> ['Private' 'Self-employed' 'Govt_job' 'children' 'Never_worked']
Residence_type -> ['Urban' 'Rural']
smoking_status -> ['formerly smoked' 'never smoked' 'smokes' 'Unknown']
stroke -> [1 0]
```


```
stroke_dataset['gender'].value_counts().plot.pie(autopct='%1.1f%%', explode=[0.2, 0.2, 0.2
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f47c9a2d940>




```
stroke_dataset['ever_married'].value_counts().plot.pie(autopct='%1.1f%%', explode=[0.2, 0.
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f47c81f5940>




```
stroke_dataset['work_type'].value_counts().plot.pie(autopct='%1.1f%%', explode=[0.2, 0.2,
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f47c81aea90>



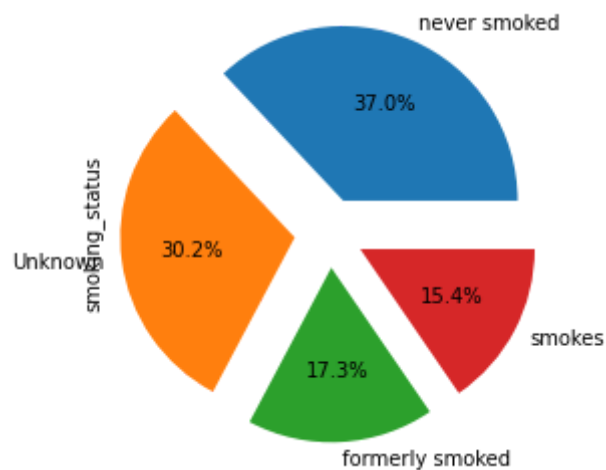
```
stroke_dataset['Residence_type'].value_counts().plot.pie(autopct='%1.1f%%', explode=[0.2,
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f47c817c670>




```
stroke_dataset['smoking_status'].value_counts().plot.pie(autopct='%1.1f%%', explode=[0.2,
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f47c80c0790>




```
stroke_dataset['hypertension'].value_counts().plot.pie(autopct='%1.1f%%', explode=[0.2, 0.
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f47c9a0fac0>
```



```
stroke_dataset['heart_disease'].value_counts().plot.pie(autopct='%1.1f%%', explode=[0.2, 0
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f47c8044dc0>
```



## IDENTIFYING THE NULL VALUES

```
stroke_dataset.isna().sum()
```

```
id                0
gender            0
age              0
hypertension      0
heart_disease     0
ever_married      0
work_type         0
Residence_type    0
avg_glucose_level 0
bmi              201
smoking_status    0
stroke            0
dtype: int64
```

## DATA PREPROCESSING

```
# Exclude id column
stroke_dataset.pop('id')
# exclude the rows containing Other and Never_worked
stroke_dataset = stroke_dataset[stroke_dataset.gender != 'Other']
stroke_dataset = stroke_dataset[stroke_dataset.work_type != 'Never_worked']
```

## REPLACING NULL VALUES


```
stroke_dataset['bmi'].fillna(stroke_dataset['bmi'].median(), inplace=True)

stroke_dataset.isna().sum()
```

```
gender          0
age             0
hypertension    0
heart_disease   0
ever_married    0
work_type       0
Residence_type  0
avg_glucose_level 0
bmi             0
smoking_status  0
stroke          0
dtype: int64
```

## DISPLAYING THE HEAT MAP

```
sns.heatmap(data=stroke_dataset.corr(), annot=True, cmap="Purples")
fig=plt.gcf()
fig.set_size_inches(15,12)
plt.show()
```



```
y_train = stroke_dataset.pop('stroke')
y_train = np.asarray(y_train).astype('float32')
```

```
print("Labels shape:", y_train.shape)
```

```
Labels shape: (5087,)
```

```
# get features
le = LabelEncoder()
ss = StandardScaler()

for col in stroke_dataset.columns:
    if stroke_dataset[col].dtype == 'object':
        stroke_dataset[col] = le.fit_transform(stroke_dataset[col])
    elif stroke_dataset[col].dtype == 'int64':
        stroke_dataset[col] = np.asarray(stroke_dataset[col]).astype('float64')
    else:
        stroke_dataset[col] = ss.fit_transform(stroke_dataset[[col]])
```

```
x_train = np.asarray(stroke_dataset)
print("Features shape:", x_train.shape)
```

```
Features shape: (5087, 10)
```

```
stroke_dataset.drop('Residence_type', axis=1, inplace=True)
```

```
sm = SMOTE()
x_train, y_train = sm.fit_resample(x_train, y_train)
part_x_train, x_test, part_y_train, y_test = train_test_split(x_train, y_train, test_size=
x_bal, y_bal = sm.fit_resample(part_x_train, part_y_train)
x_test, x_val, y_test, y_val = train_test_split(x_test, y_test, test_size=0.5, random_stat
```

## BUILDING THE MODEL

```
model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(x_bal.shape[1],)))
model.add(layers.Dropout(0.2))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.2))
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dropout(0.2))
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dropout(0.2))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.2))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.2))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.2))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.2))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dropout(0.2))
model.add(layers.Dense(32, activation='relu'))
model.add(layers.Dropout(0.2))
model.add(layers.Dense(1, activation='sigmoid'))
early_stopping = EarlyStopping(monitor='val_loss', min_delta=0.01, patience=20, restore_best_weights=True)
model.compile(optimizer=optimizers.Adam(learning_rate=2e-4), loss='binary_crossentropy', metrics=['accuracy'])
history = model.fit(x_bal, y_bal, epochs=79, batch_size=128, validation_data=(x_val, y_val))
```

```

Epoch 67/79
53/53 [=====] - 1s 10ms/step - loss: 0.3380 - accuracy:
Epoch 68/79
53/53 [=====] - 1s 10ms/step - loss: 0.3313 - accuracy:
Epoch 69/79
53/53 [=====] - 1s 9ms/step - loss: 0.3313 - accuracy: 0
Epoch 70/79
53/53 [=====] - 1s 11ms/step - loss: 0.3223 - accuracy:
Epoch 71/79
53/53 [=====] - 1s 11ms/step - loss: 0.3279 - accuracy:
Epoch 72/79
53/53 [=====] - 1s 10ms/step - loss: 0.3241 - accuracy:
Epoch 73/79
53/53 [=====] - 1s 10ms/step - loss: 0.3272 - accuracy:
Epoch 74/79
53/53 [=====] - 1s 10ms/step - loss: 0.3229 - accuracy:
Epoch 75/79
53/53 [=====] - 1s 10ms/step - loss: 0.3204 - accuracy:
Epoch 76/79
53/53 [=====] - 1s 10ms/step - loss: 0.3206 - accuracy:
Epoch 77/79
53/53 [=====] - 1s 10ms/step - loss: 0.3167 - accuracy:
Epoch 78/79
53/53 [=====] - 1s 10ms/step - loss: 0.3224 - accuracy:
Epoch 79/79
53/53 [=====] - 1s 10ms/step - loss: 0.3224 - accuracy:

```

```
model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	704
dropout (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 128)	8320
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 256)	33024
dropout_2 (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 256)	65792
dropout_3 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 128)	32896
dropout_4 (Dropout)	(None, 128)	0
dense_5 (Dense)	(None, 128)	16512
dropout_5 (Dropout)	(None, 128)	0
dense_6 (Dense)	(None, 128)	16512




dropout_6 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 128)	16512
dropout_7 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 64)	8256
dropout_8 (Dropout)	(None, 64)	0
dense_9 (Dense)	(None, 32)	2080
dropout_9 (Dropout)	(None, 32)	0
dense_10 (Dense)	(None, 1)	33

```
=====
Total params: 200,641
Trainable params: 200,641
Non-trainable params: 0
=====
```


## TRAINING AND VALIDATION LOSS

```
history= history.history
loss_values = history['loss']
val_loss_values = history['val_loss']
epochs = range(1, len(history['accuracy']) + 1)
plt.plot(epochs, loss_values, 'r', label='Training loss')
plt.plot(epochs, val_loss_values, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



## TARINING AND VALIDATION ACCURACY

```
plt.clf()
acc_values = history['accuracy']
val_acc_values = history['val_accuracy']
plt.plot(epochs, history['accuracy'], 'r', label='Training acc')
plt.plot(epochs, history['val_accuracy'], 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



```
y_pred = model.predict(x_test)
y_pred = [1.0 if p > 0.5 else 0 for p in y_pred]
```


46/46 [=====] - 0s 2ms/step

y\_pred

```
[0,
 1.0,
 1.0,
 1.0,
 1.0,
 1.0,
 0,
 1.0,
 1.0,
 1.0,
 0,
 1.0,
 0,
 1.0,
 1.0,
 0,
 1.0,
 1.0,
 1.0,
 0,
 1.0,
 0,
 1.0,
 0]
```

1.0,  
1.0,  
1.0,  
1.0,  
0,  
1.0,  
0,  
1.0,  
1.0,  
1.0,  
0,  
0,  
0,  
1.0,  
1.0,  
0,  
1.0,  
1.0,  
0,  
0,  
0,  
1.0,  
0,  
0,  
0,  
0,  
0,  
0,  
1.0,  
0,  
1.0,  
0,  
0,  
1.0,  
1.0,  
0,  
0,  
1.0,  
1.0,  
0,

```
plt.figure(figsize=(15,6))  
sns.countplot(x=y_pred, palette='Blues')  
plt.xticks(rotation = 90)  
plt.show()
```



```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0.0	0.98	0.73	0.84	715
1.0	0.79	0.99	0.88	736
accuracy			0.86	1451
macro avg	0.89	0.86	0.86	1451
weighted avg	0.88	0.86	0.86	1451